

Simple Linear Regression

```
import pandas as pd
import numpy as np
df = pd.read_csv('salary_data.csv')
```

```
df
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

```
df.head()
```

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0

```
df.shape
```

```
(30, 2)
```

```
df[['YearsExperience']] #independetn variable
```

	YearsExperience
0	1.1
1	1.3
2	1.5
3	2.0
4	2.2
5	2.9
6	3.0
7	3.2
8	3.2
9	3.7
10	3.9
11	4.0
12	4.0
13	4.1
14	4.5
15	4.9
16	5.1
17	5.3
18	5.9
19	6.0
20	6.8
21	7.1
22	7.9
23	8.2
24	8.7
25	9.0
26	9.5
27	9.6
28	10.3
29	10.5

```
df[['Salary']].head()
```

	Salary
0	39343.0
1	46205.0
2	37731.0
3	43525.0
4	39891.0

```
df.describe()
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785

```
min          1.100000    37731.000000
25%          3.200000    56720.750000
50%          4.700000    65237.000000
75%          7.700000   100544.750000
max          10.500000   122391.000000
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 30 entries, 0 to 29
```

```
Data columns (total 2 columns):
```

#	Column	Non-Null Count	Dtype
0	YearsExperience	30 non-null	float64
1	Salary	30 non-null	float64

```
dtypes: float64(2)
```

```
memory usage: 608.0 bytes
```

```
df.isnull().sum()
```

```
YearsExperience    0
```


```
Salary            0
```

```
dtype: int64
```

```
import matplotlib.pyplot as plt #data visualization
```

```
plt.scatter(df.YearsExperience,df.Salary) # scatter plot
```

```
<matplotlib.collections.PathCollection at 0x1ec388414c0>
```



Independent and Dependent variables

```
x = df.iloc[:,0:1]
```

x

	YearsExperience
0	1.1
1	1.3
2	1.5
3	2.0
4	2.2
5	2.9
6	3.0
7	3.2
8	3.2
9	3.7
10	3.9
11	4.0
12	4.0
13	4.1
14	4.5
15	4.9
16	5.1
17	5.3
18	5.9
19	6.0
20	6.8
21	7.1
22	7.9
23	8.2
24	8.7
25	9.0
26	9.5
27	9.6
28	10.3
29	10.5

```
y = df.iloc[:,1:]
```

y

	Salary
0	39343.0
1	46205.0
2	37731.0
3	43525.0
4	39891.0
5	56642.0
6	60150.0
7	54445.0
8	64445.0
9	57189.0

```
10    63218.0
11    55794.0
12    56957.0
13    57081.0
14    61111.0
15    67938.0
16    66029.0
17    83088.0
18    81363.0
19    93940.0
20    91738.0
21    98273.0
22   101302.0
23   113812.0
24   109431.0
25   105582.0
26   116969.0
27   112635.0
28   122391.0
29   121872.0
```

```
df.shape
```

```
(30, 2)
```

Train, Test & Split

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test =
train_test_split(x,y,test_size=0.2,random_state=0)
```

```
x_train.shape
```

```
(24, 1)
```

```
x_test.shape
```

```
(6, 1)
```

```
y_train.shape
```

```
(24, 1)
```

```
y_test.shape
```

```
(6, 1)
```

Model Building

```
from sklearn.linear_model import LinearRegression
sl = LinearRegression()
```

```
sl.fit(x_train,y_train) # simple linear regression is created with training data
```

```
LinearRegression()
```

```
x_test
```

	YearsExperience
2	1.5
28	10.3
13	4.1
10	3.9
26	9.5
24	8.7

```
y_test
```

	Salary
2	37731.0
28	122391.0
13	57081.0
10	63218.0
26	116969.0
24	109431.0

```
sl.predict(x_test)
```

```
array([[ 40748.96184072],  
       [122699.62295594],  
       [ 64961.65717022],  
       [ 63099.14214487],  
       [115249.56285456],  
       [107799.50275317]])
```

```
from sklearn.metrics import r2_score  
r2_score(sl.predict(x_test),y_test)
```

```
0.986482673117654
```

```
sl.predict([[9.3]]) # testing with unseen data
```


```
array([[113387.04782921]])
```

```
sl.predict([[12.3]]) # 12.3 is experience, so we are getting salary
```

```
array([[141324.7732094]])
```

```
plt.scatter(x_train,y_train)  
plt.plot(x_train,sl.predict(x_train),'r')
```

```
[<matplotlib.lines.Line2D at 0x1ec38d1a3d0>]
```



3 Multiple Linear Regression

```
import numpy as np
import pandas as pd
```

```
df = pd.read_csv('50_Startups (4).csv')
```

```
df
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94
...
103	119943.24	156547.42	256512.92	Florida	132602.65
104	114523.61	122616.84	261776.23	New York	129917.04
105	78013.11	121597.55	264346.06	California	126992.93
106	94657.16	145077.58	282574.31	New York	125370.37
107	91749.16	114175.79	294919.57	Florida	124266.90

```
[108 rows x 5 columns]
```

```
df.head()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39

3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
df.shape
```

```
(108, 5)
```

```
df.tail()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
103	119943.24	156547.42	256512.92	Florida	132602.65
104	114523.61	122616.84	261776.23	New York	129917.04
105	78013.11	121597.55	264346.06	California	126992.93
106	94657.16	145077.58	282574.31	New York	125370.37
107	91749.16	114175.79	294919.57	Florida	124266.90

```
df.isnull().sum()
```

```
R&D Spend      0
Administration  0
Marketing Spend  0
State           0
Profit          0
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 108 entries, 0 to 107
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	R&D Spend	108 non-null	float64
1	Administration	108 non-null	float64
2	Marketing Spend	108 non-null	float64
3	State	108 non-null	object
4	Profit	108 non-null	float64

```
dtypes: float64(4), object(1)
```

```
memory usage: 4.3+ KB
```

```
df.describe() # to get statical information
```

	R&D Spend	Administration	Marketing Spend	Profit
count	108.000000	108.000000	108.000000	108.000000
mean	74959.338704	121750.788889	214952.664722	113523.760000
std	44996.368152	27322.385654	117937.942120	38991.013654
min	0.000000	51283.140000	0.000000	14681.400000
25%	38558.510000	105077.645000	134050.070000	90708.190000
50%	75791.365000	122699.795000	239452.750000	109543.120000
75%	101913.080000	145077.580000	298664.470000	141585.520000
max	165349.200000	182645.560000	471784.100000	192261.830000

```
df.State.value_counts()
```



```
New York      39
California    36
Florida       33
Name: State, dtype: int64
```

```
df.columns
```

```
Index(['R&D Spend', 'Administration', 'Marketing Spend', 'State',
      'Profit'], dtype='object')
```

Label Encoder

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
df.State = le.fit_transform(df.State)
df
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	2	192261.83
1	162597.70	151377.59	443898.53	0	191792.06
2	153441.51	101145.55	407934.54	1	191050.39
3	144372.41	118671.85	383199.62	2	182901.99
4	142107.34	91391.77	366168.42	1	166187.94
...
103	119943.24	156547.42	256512.92	1	132602.65
104	114523.61	122616.84	261776.23	2	129917.04
105	78013.11	121597.55	264346.06	0	126992.93
106	94657.16	145077.58	282574.31	2	125370.37
107	91749.16	114175.79	294919.57	1	124266.90

```
[108 rows x 5 columns]
```

```
df.State.value_counts() # encoded
```

```
2      39
0      36
1      33
Name: State, dtype: int64
```

```
df.head()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	2	192261.83
1	162597.70	151377.59	443898.53	0	191792.06
2	153441.51	101145.55	407934.54	1	191050.39
3	144372.41	118671.85	383199.62	2	182901.99
4	142107.34	91391.77	366168.42	1	166187.94

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 108 entries, 0 to 107
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   R&D Spend              108 non-null   float64
1   Administration         108 non-null   float64
2   Marketing Spend        108 non-null   float64
3   State                  108 non-null   int32
4   Profit                 108 non-null   float64
dtypes: float64(4), int32(1)
memory usage: 3.9 KB

```

Independent and Dependent variables

```

x = df.iloc[:,4]
x

```

	R&D Spend	Administration	Marketing Spend	State
0	165349.20	136897.80	471784.10	2
1	162597.70	151377.59	443898.53	0
2	153441.51	101145.55	407934.54	1
3	144372.41	118671.85	383199.62	2
4	142107.34	91391.77	366168.42	1
...
103	119943.24	156547.42	256512.92	1
104	114523.61	122616.84	261776.23	2
105	78013.11	121597.55	264346.06	0
106	94657.16	145077.58	282574.31	2
107	91749.16	114175.79	294919.57	1

```

[108 rows x 4 columns]

```

```

y = df.iloc[:,4:]
y

```

	Profit
0	192261.83
1	191792.06
2	191050.39
3	182901.99
4	166187.94
...	...
103	132602.65
104	129917.04
105	126992.93
106	125370.37
107	124266.90

```
[108 rows x 1 columns]
```

Train, Test & Split

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size =
0.2,random_state=0)
```

```
df.shape
```

```
(108, 5)
```

```
x_train.shape#remaining all the variables
```

```
(86, 4)
```

```
y_train.shape #profit
```

```
(86, 1)
```

```
x_test.shape
```

```
(22, 4)
```

```
from sklearn.linear_model import LinearRegression
slr = LinearRegression()
slr.fit(x_train,y_train)
```

```
LinearRegression()
```

```
x_test.head()
```

	R&D Spend	Administration	Marketing Spend	State
84	1000.23	124153.04	1903.93	2
10	101913.08	110594.11	229160.95	1
75	28663.76	127056.21	201126.82	1
2	153441.51	101145.55	407934.54	1
24	77044.01	99281.34	140574.81	2

```
y_test.head()
```

	Profit
84	64926.08
10	146121.95
75	90708.19
2	191050.39
24	108552.04

```
x_test[0:5]
```

	R&D Spend	Administration	Marketing Spend	State
84	1000.23	124153.04	1903.93	2

10	101913.08	110594.11	229160.95	1
75	28663.76	127056.21	201126.82	1
2	153441.51	101145.55	407934.54	1
24	77044.01	99281.34	140574.81	2

```
mlr_pred = slr.predict(x_test)
mlr_pred[0:5]
```

```
array([[ 48379.24868384],
       [134848.91924675],
       [ 76483.10965219],
       [181561.78529195],
       [112966.00035119]])
```

```
slr.predict(x_test[0:5])
```

```
array([[ 48379.24868384],
       [134848.91924675],
       [ 76483.10965219],
       [181561.78529195],
       [112966.00035119]])
```

```
import numpy as np
from sklearn.metrics import r2_score
r2_score(slr.predict(x_test), y_test)
```

```
0.9314720388994493
```

```
slr.predict([[123120.54, 12321.32, 567800.34, 0]])
```

```
array([[163972.11719139]])
```

Polynomial Regression

```
df = pd.read_csv('salaries_data.csv')
```

```
df
```

	Position	Level	Salary
0	Business Analyst	1	45000
1	Junior Consultant	2	50000
2	Senior Consultant	3	60000
3	Manager	4	80000
4	Country Manager	5	110000
5	Region Manager	6	150000
6	Partner	7	200000
7	Senior Partner	8	300000
8	C-level	9	500000
9	CEO	10	1000000

```
df.head()
```

	Position	Level	Salary
0	Business Analyst	1	45000
1	Junior Consultant	2	50000
2	Senior Consultant	3	60000
3	Manager	4	80000
4	Country Manager	5	110000

```
df.corr()
```

	Level	Salary
Level	1.000000	0.817949
Salary	0.817949	1.000000

Independent and Dependent variables

```
x = df.iloc[:,1:2]
```

```
x
```

	Level
0	1
1	2
2	3
3	4
4	5
5	6
6	7
7	8
8	9
9	10


```
y = df.iloc[:,2:]
```

```
y
```

	Salary
0	45000
1	50000
2	60000
3	80000
4	110000
5	150000
6	200000
7	300000
8	500000
9	1000000

```
import matplotlib.pyplot as plt
plt.scatter(x,y)
```

```
<matplotlib.collections.PathCollection at 0x1ec38dale80>
```



```
from sklearn.preprocessing import PolynomialFeatures
pr = PolynomialFeatures(degree = 3)
```

```
xp = pr.fit_transform(x)
```

```
xp
```


```
array([[ 1.,  1.,  1.,  1.],
       [ 1.,  2.,  4.,  8.],
       [ 1.,  3.,  9., 27.],
       [ 1.,  4., 16., 64.],
       [ 1.,  5., 25., 125.],
       [ 1.,  6., 36., 216.],
       [ 1.,  7., 49., 343.],
       [ 1.,  8., 64., 512.],
       [ 1.,  9., 81., 729.],
       [ 1., 10., 100., 1000.]])
```

```
lr = LinearRegression()
lr.fit(xp,y)
```

```
LinearRegression()
```

```
plt.scatter(x,y)
plt.plot(x,lr.predict(xp),'r')
```

```
[<matplotlib.lines.Line2D at 0x1ec38e02fa0>]
```



Decision Tree Regression

- # It build the regression in tree structure
- # which has decision and leaf nodes
- # it split the big thing into smaller ones
- # applicable for both classification and regression problems
- # it had root node and internal decision nodes
- # and it has subtress and children nodes and we can also perform cloning(removing of the unwanted nodes from the tree)
- # we can best attribute

```
df = pd.read_csv('50_Startups (4).csv')
```

```
df.head()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

Model Building

```
from sklearn.tree import DecisionTreeRegressor
dt = DecisionTreeRegressor()
dt.fit(x_train,y_train)
```

```
DecisionTreeRegressor()
```

```
dt.predict(x_test)
array([ 42559.73, 146121.95,  90708.19, 191050.39, 108552.04, 144259.4
,
      124266.9 , 155752.6 , 126992.93,  42559.73, 101004.64,
110352.25,
      42559.73, 111313.02,  89949.14, 134307.35, 134307.35,  96712.8
,
      49490.75, 129917.04, 132602.65, 152211.77])
y_test
```

	Profit
84	64926.08
10	146121.95
75	90708.19
2	191050.39
24	108552.04
100	144259.40
107	124266.90
7	155752.60
16	126992.93
86	42559.73
68	101004.64
22	110352.25
45	64926.08
60	111313.02
76	89949.14
52	134307.35
13	134307.35
73	96712.80
85	49490.75
54	129917.04
103	132602.65
8	152211.77

Random forest Regression

```
# if we have more decision tress in the dataset and most of them are
supporting categorical data
# assemble uses bagging and boosting
# bagging output is majority good
# boosting creates the accurate model
# overcome the over fit problem
```

```
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(n_estimators = 5,random_state = 0)
rf.fit(x_train,y_train)
```

C:\Users\bharg\AppData\Local\Temp\ipykernel_1992\921291185.py:3:
DataConversionWarning: A column-vector y was passed when a 1d array

was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
rf.fit(x_train,y_train)
```

```
RandomForestRegressor(n_estimators=5, random_state=0)
```

```
rfr = rf.predict(x_test) # predicted data
rfr
```

```
array([ 36993.004, 145749.44 ,  84934.656, 187791.03 , 128036.856,
        144259.4  , 124812.106, 155044.434, 126992.93 ,  25832.732,
        101460.188, 110352.25 ,  36993.004, 115556.692,  92681.052,
        132844.466, 132844.466,  96680.32 ,  32803.542, 133513.412,
        136034.112, 131440.084])
```

```
y_test # actual data
```

	Profit
84	64926.08
10	146121.95
75	90708.19
2	191050.39
24	108552.04
100	144259.40
107	124266.90
7	155752.60
16	126992.93
86	42559.73
68	101004.64
22	110352.25
45	64926.08
60	111313.02
76	89949.14
52	134307.35
13	134307.35
73	96712.80
85	49490.75
54	129917.04
103	132602.65
8	152211.77

```
rfscore = r2_score(rfr,y_test)
```

```
rfscore
```

```
0.922343446522065
```

Classification Problems

```
# we have multi class classifier
```

```
#1.garbage (glass,metal,fiber,plastic)
```

```
#2.more than two we contain
# binary classifier
#1.two possible outcomes(no/yes and reject/approve)
```

Decision tree classifier

```
df = pd.read_csv('loan_prediction.csv')
```

```
df.head()
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 614 entries, 0 to 613
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	Loan_ID	614 non-null	object
1	Gender	601 non-null	object
2	Married	611 non-null	object
3	Dependents	599 non-null	object
4	Education	614 non-null	object
5	Self_Employed	582 non-null	object
6	ApplicantIncome	614 non-null	int64
7	CoapplicantIncome	614 non-null	float64
8	LoanAmount	592 non-null	float64
9	Loan_Amount_Term	600 non-null	float64
10	Credit_History	564 non-null	float64
11	Property_Area	614 non-null	object

```
12 Loan_Status      614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

```
df.isnull().sum()
```

```
Loan_ID      0
Gender       13
Married       3
Dependents   15
Education     0
Self_Employed 32
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount   22
Loan_Amount_Term 14
Credit_History 50
Property_Area 0
Loan_Status   0
dtype: int64
```

```
df['Gender'].isnull().sum()
```

```
13
```

```
df['Gender'].unique()
```

```
array(['Male', 'Female', nan], dtype=object)
```

```
df['Married'].unique()
```

```
array(['No', 'Yes', nan], dtype=object)
```

```
df['Dependents'].unique()
```

```
array(['0', '1', '2', '3+', nan], dtype=object)
```

```
df['Gender'] = df['Gender'].map({'Male':1,'Female':0})
```

```
df['Married'] = df['Married'].map({'Yes':1,'No':0})
```

```
#df['Gender'].unique()
```

```
df['Married'].unique()
```

```
array([ 0.,  1., nan])
```

Logistic Regression

```
#credit problems all are resolve by using the logistic regression
#ln(p/1-p)
```

```
import pandas as pd
import numpy as np #scientific calculations
df = pd.read_csv('Social_Network_Ads.csv')
```

```
df.head()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
df.shape
```

```
(400, 5)
```

```
df.describe()
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID                400 non-null   int64
1   Gender                 400 non-null   object
2   Age                    400 non-null   int64
3   EstimatedSalary        400 non-null   int64
4   Purchased              400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

```
df.isnull().sum()
```

```
User ID      0
Gender       0
Age          0
EstimatedSalary  0
Purchased    0
dtype: int64
```

```
df['Purchased'].unique()
```

```
array([0, 1], dtype=int64)
```

Independent and Dependent Variables

```
x =df.iloc[:,[2,3]].values
```

x

```
array([[ 19, 19000],
       [ 35, 20000],
       [ 26, 43000],
       [ 27, 57000],
       [ 19, 76000],
       [ 27, 58000],
       [ 27, 84000],
       [ 32, 150000],
       [ 25, 33000],
       [ 35, 65000],
       [ 26, 80000],
       [ 26, 52000],
       [ 20, 86000],
       [ 32, 18000],
       [ 18, 82000],
       [ 29, 80000],
       [ 47, 25000],
       [ 45, 26000],
       [ 46, 28000],
       [ 48, 29000],
       [ 45, 22000],
       [ 47, 49000],
       [ 48, 41000],
       [ 45, 22000],
       [ 46, 23000],
       [ 47, 20000],
       [ 49, 28000],
       [ 47, 30000],
       [ 29, 43000],
       [ 31, 18000],
       [ 31, 74000],
       [ 27, 137000],
       [ 21, 16000],
       [ 28, 44000],
       [ 27, 90000],
       [ 35, 27000],
       [ 33, 28000],
       [ 30, 49000],
       [ 26, 72000],
       [ 27, 31000],
       [ 27, 17000],
       [ 33, 51000],
       [ 35, 108000],
       [ 30, 15000],
       [ 28, 84000],
```

```
[ 23, 20000],
[ 25, 79000],
[ 27, 54000],
[ 30, 135000],
[ 31, 89000],
[ 24, 32000],
[ 18, 44000],
[ 29, 83000],
[ 35, 23000],
[ 27, 58000],
[ 24, 55000],
[ 23, 48000],
[ 28, 79000],
[ 22, 18000],
[ 32, 117000],
[ 27, 20000],
[ 25, 87000],
[ 23, 66000],
[ 32, 120000],
[ 59, 83000],
[ 24, 58000],
[ 24, 19000],
[ 23, 82000],
[ 22, 63000],
[ 31, 68000],
[ 25, 80000],
[ 24, 27000],
[ 20, 23000],
[ 33, 113000],
[ 32, 18000],
[ 34, 112000],
[ 18, 52000],
[ 22, 27000],
[ 28, 87000],
[ 26, 17000],
[ 30, 80000],
[ 39, 42000],
[ 20, 49000],
[ 35, 88000],
[ 30, 62000],
[ 31, 118000],
[ 24, 55000],
[ 28, 85000],
[ 26, 81000],
[ 35, 50000],
[ 22, 81000],
[ 30, 116000],
[ 26, 15000],
[ 29, 28000],
[ 29, 83000],
```

[35, 44000],
[35, 25000],
[28, 123000],
[35, 73000],
[28, 37000],
[27, 88000],
[28, 59000],
[32, 86000],
[33, 149000],
[19, 21000],
[21, 72000],
[26, 35000],
[27, 89000],
[26, 86000],
[38, 80000],
[39, 71000],
[37, 71000],
[38, 61000],
[37, 55000],
[42, 80000],
[40, 57000],
[35, 75000],
[36, 52000],
[40, 59000],
[41, 59000],
[36, 75000],
[37, 72000],
[40, 75000],
[35, 53000],
[41, 51000],
[39, 61000],
[42, 65000],
[26, 32000],
[30, 17000],
[26, 84000],
[31, 58000],
[33, 31000],
[30, 87000],
[21, 68000],
[28, 55000],
[23, 63000],
[20, 82000],
[30, 107000],
[28, 59000],
[19, 25000],
[19, 85000],
[18, 68000],
[35, 59000],
[30, 89000],
[34, 25000],

[24, 89000],
[27, 96000],
[41, 30000],
[29, 61000],
[20, 74000],
[26, 15000],
[41, 45000],
[31, 76000],
[36, 50000],
[40, 47000],
[31, 15000],
[46, 59000],
[29, 75000],
[26, 30000],
[32, 135000],
[32, 100000],
[25, 90000],
[37, 33000],
[35, 38000],
[33, 69000],
[18, 86000],
[22, 55000],
[35, 71000],
[29, 148000],
[29, 47000],
[21, 88000],
[34, 115000],
[26, 118000],
[34, 43000],
[34, 72000],
[23, 28000],
[35, 47000],
[25, 22000],
[24, 23000],
[31, 34000],
[26, 16000],
[31, 71000],
[32, 117000],
[33, 43000],
[33, 60000],
[31, 66000],
[20, 82000],
[33, 41000],
[35, 72000],
[28, 32000],
[24, 84000],
[19, 26000],
[29, 43000],
[19, 70000],
[28, 89000],

[34, 43000],
[30, 79000],
[20, 36000],
[26, 80000],
[35, 22000],
[35, 39000],
[49, 74000],
[39, 134000],
[41, 71000],
[58, 101000],
[47, 47000],
[55, 130000],
[52, 114000],
[40, 142000],
[46, 22000],
[48, 96000],
[52, 150000],
[59, 42000],
[35, 58000],
[47, 43000],
[60, 108000],
[49, 65000],
[40, 78000],
[46, 96000],
[59, 143000],
[41, 80000],
[35, 91000],
[37, 144000],
[60, 102000],
[35, 60000],
[37, 53000],
[36, 126000],
[56, 133000],
[40, 72000],
[42, 80000],
[35, 147000],
[39, 42000],
[40, 107000],
[49, 86000],
[38, 112000],
[46, 79000],
[40, 57000],
[37, 80000],
[46, 82000],
[53, 143000],
[42, 149000],
[38, 59000],
[50, 88000],
[56, 104000],
[41, 72000],

[51, 146000],
[35, 50000],
[57, 122000],
[41, 52000],
[35, 97000],
[44, 39000],
[37, 52000],
[48, 134000],
[37, 146000],
[50, 44000],
[52, 90000],
[41, 72000],
[40, 57000],
[58, 95000],
[45, 131000],
[35, 77000],
[36, 144000],
[55, 125000],
[35, 72000],
[48, 90000],
[42, 108000],
[40, 75000],
[37, 74000],
[47, 144000],
[40, 61000],
[43, 133000],
[59, 76000],
[60, 42000],
[39, 106000],
[57, 26000],
[57, 74000],
[38, 71000],
[49, 88000],
[52, 38000],
[50, 36000],
[59, 88000],
[35, 61000],
[37, 70000],
[52, 21000],
[48, 141000],
[37, 93000],
[37, 62000],
[48, 138000],
[41, 79000],
[37, 78000],
[39, 134000],
[49, 89000],
[55, 39000],
[37, 77000],
[35, 57000],

[36, 63000],
[42, 73000],
[43, 112000],
[45, 79000],
[46, 117000],
[58, 38000],
[48, 74000],
[37, 137000],
[37, 79000],
[40, 60000],
[42, 54000],
[51, 134000],
[47, 113000],
[36, 125000],
[38, 50000],
[42, 70000],
[39, 96000],
[38, 50000],
[49, 141000],
[39, 79000],
[39, 75000],
[54, 104000],
[35, 55000],
[45, 32000],
[36, 60000],
[52, 138000],
[53, 82000],
[41, 52000],
[48, 30000],
[48, 131000],
[41, 60000],
[41, 72000],
[42, 75000],
[36, 118000],
[47, 107000],
[38, 51000],
[48, 119000],
[42, 65000],
[40, 65000],
[57, 60000],
[36, 54000],
[58, 144000],
[35, 79000],
[38, 55000],
[39, 122000],
[53, 104000],
[35, 75000],
[38, 65000],
[47, 51000],
[47, 105000],

[41, 63000],
[53, 72000],
[54, 108000],
[39, 77000],
[38, 61000],
[38, 113000],
[37, 75000],
[42, 90000],
[37, 57000],
[36, 99000],
[60, 34000],
[54, 70000],
[41, 72000],
[40, 71000],
[42, 54000],
[43, 129000],
[53, 34000],
[47, 50000],
[42, 79000],
[42, 104000],
[59, 29000],
[58, 47000],
[46, 88000],
[38, 71000],
[54, 26000],
[60, 46000],
[60, 83000],
[39, 73000],
[59, 130000],
[37, 80000],
[46, 32000],
[46, 74000],
[42, 53000],
[41, 87000],
[58, 23000],
[42, 64000],
[48, 33000],
[44, 139000],
[49, 28000],
[57, 33000],
[56, 60000],
[49, 39000],
[39, 71000],
[47, 34000],
[48, 35000],
[48, 33000],
[47, 23000],
[45, 45000],
[60, 42000],
[39, 59000],

```

[    46, 41000],
[    51, 23000],
[    50, 20000],
[    36, 33000],
[    49, 36000]], dtype=int64)

y = df.iloc[:,4].values

y
array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,
1,
      1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0,
1,
      0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1,
0,
      1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1,
0,
      1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0,
1,
      0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0,
1,
      1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1,
1,
      0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1,
0,
      1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0,
1,
      0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0,
1,
      1, 1, 0, 1], dtype=int64)

df['Purchased'].value_counts()

```

```
0    257
1    143
Name: Purchased, dtype: int64
```

Train,Test and split

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size =
0.2,random_state = 0)
```

feature scaling

#cannot apply for regression problem, we can only implement for classification problems
$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$ if num == den then $x' = 1$ it is a maximum case in feature scaling
#Standardization is one of the scaling technique
$x' = \frac{x - \mu}{\sigma}$
#nomalization is useful when the distribution is not in the form of guassian distribution
1.x' = 0 min
2.x' = 1 max
3.x' =inbetween 0 and 1
#standardization is not useful when the distribution is not in the form of guassian distribution

File

"C:\Users\bharg\AppData\Local\Temp\ipykernel_1992\220819968.py", line 6

```
1.x' = 0 min
   ^
```

SyntaxError: invalid syntax

```
from sklearn.preprocessing import StandardScaler
st = StandardScaler()
x_train= st.fit_transform(x_train)

x_train
x_test[0:5]
x_test = st.transform(x_test)
x_test[0:5]
```

Model Building

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(x_train,y_train)

LogisticRegression()

lrpred = lr.predict(x_test)

lrpred
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)

y_test
array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
1,
      0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0,
      1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0,
1,
      0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1], dtype=int64)
```

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(lrpred,y_test)

cm
array([[58, 22],
      [ 0,  0]], dtype=int64)
```

Accuracy score

```
from sklearn.metrics import accuracy_score
acc = accuracy_score(lrpred,y_test)

acc
0.725
```

KNN - K Nearest Neighbours

In order to calculate the distance between two points we are using different methods

1. Eculidean Distance
2. Manhattan Distance
3. Minkowski Distance

step1:

Choose the number of k of neighbours

step2:

Take the k nearest neighbours of the new datapoint, according to the euclidan distance

step3:

Among the k neighbours, count the number of datapoints **in** each category

step4:

Assign the new data points to the category where you counted the most neighbors

Model **is** ready

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
```

```
knn.fit(x_train, y_train)
```

```
knn
```

```
knnpred = knn.predict(x_test)
```

```
knnpred
```

```
cm = confusion_matrix(knnpred, y_test)
```

```
cm
```

```
from sklearn.metrics import accuracy_score
acc = accuracy_score(knnpred, y_test)
```

```
acc
```

```
knn.predict([[19, 38900]])
```

Naive Bayes

naive is specifies the occurence of one feature independent to the occurence of another feature

bayes is used to determine the hypothesis on prior knowledge

```
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(x_train, y_train)
```



```

GaussianNB()
gnbpred = gnb.predict(x_test)
gnbpred
array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
1,
      0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0,
      1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
1,
      0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1], dtype=int64)
y_test
array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
1,
      0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0,
      1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0,
1,
      0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1], dtype=int64)
cm = confusion_matrix(gnbpred,y_test)
cm
array([[56,  4],
      [ 2, 18]], dtype=int64)
gnbacc = accuracy_score(gnbpred,y_test)
gnbacc
0.925

```